



Puppet Fundamentals

This course is intended for organisations already using Puppet or Systems Administrators looking to implement Puppet as a new solution.

The course is delivered over 3 days and involves a mixture of both presentations and exercises. On completion you will be able to deploy a basic system configuration for Puppet in a Master-Client environment.

Introduction: Who We Are

Your Trainer is: Andy Singleton

A Linux DevOps trainer and consultant specialising in Linux administration, Configuration Management, Docker and OpenStack.

Around 20 years in IT, with professional certifications including

- Red Hat RHCSA, RHCE, RHCDS, and RHCA
- PuppetLabs Puppet Certification

About NobleProg

NobleProg are your training and consultancy provider for areas including Management, IT, Statistics, Programming, and Artificial Intelligence

Puppet Fundamentals

Day 1	Day 2	Day 3
Introduction	Module Structure	General Architecture
About Puppet	PuppetForge	Multi-master Puppet
Agent and Master	Classifying nodes	Foreman
Community and Enterprise	Hiera	Cobbler
Reporting	Git	Best Practise
Resources	Module Development	Case Study
Resource Relationships	Module Testing	
Variables	Deploying Code	
Facter	Master vs Masterless	
Conditionals		
Templating		
Defined Types		

Introduction: **Training Environment**

User instance

*1 CPU, 523 MB Ram, 8GB Storage
Red Hat / CentOS or Ubuntu*

SSH access: Port 22

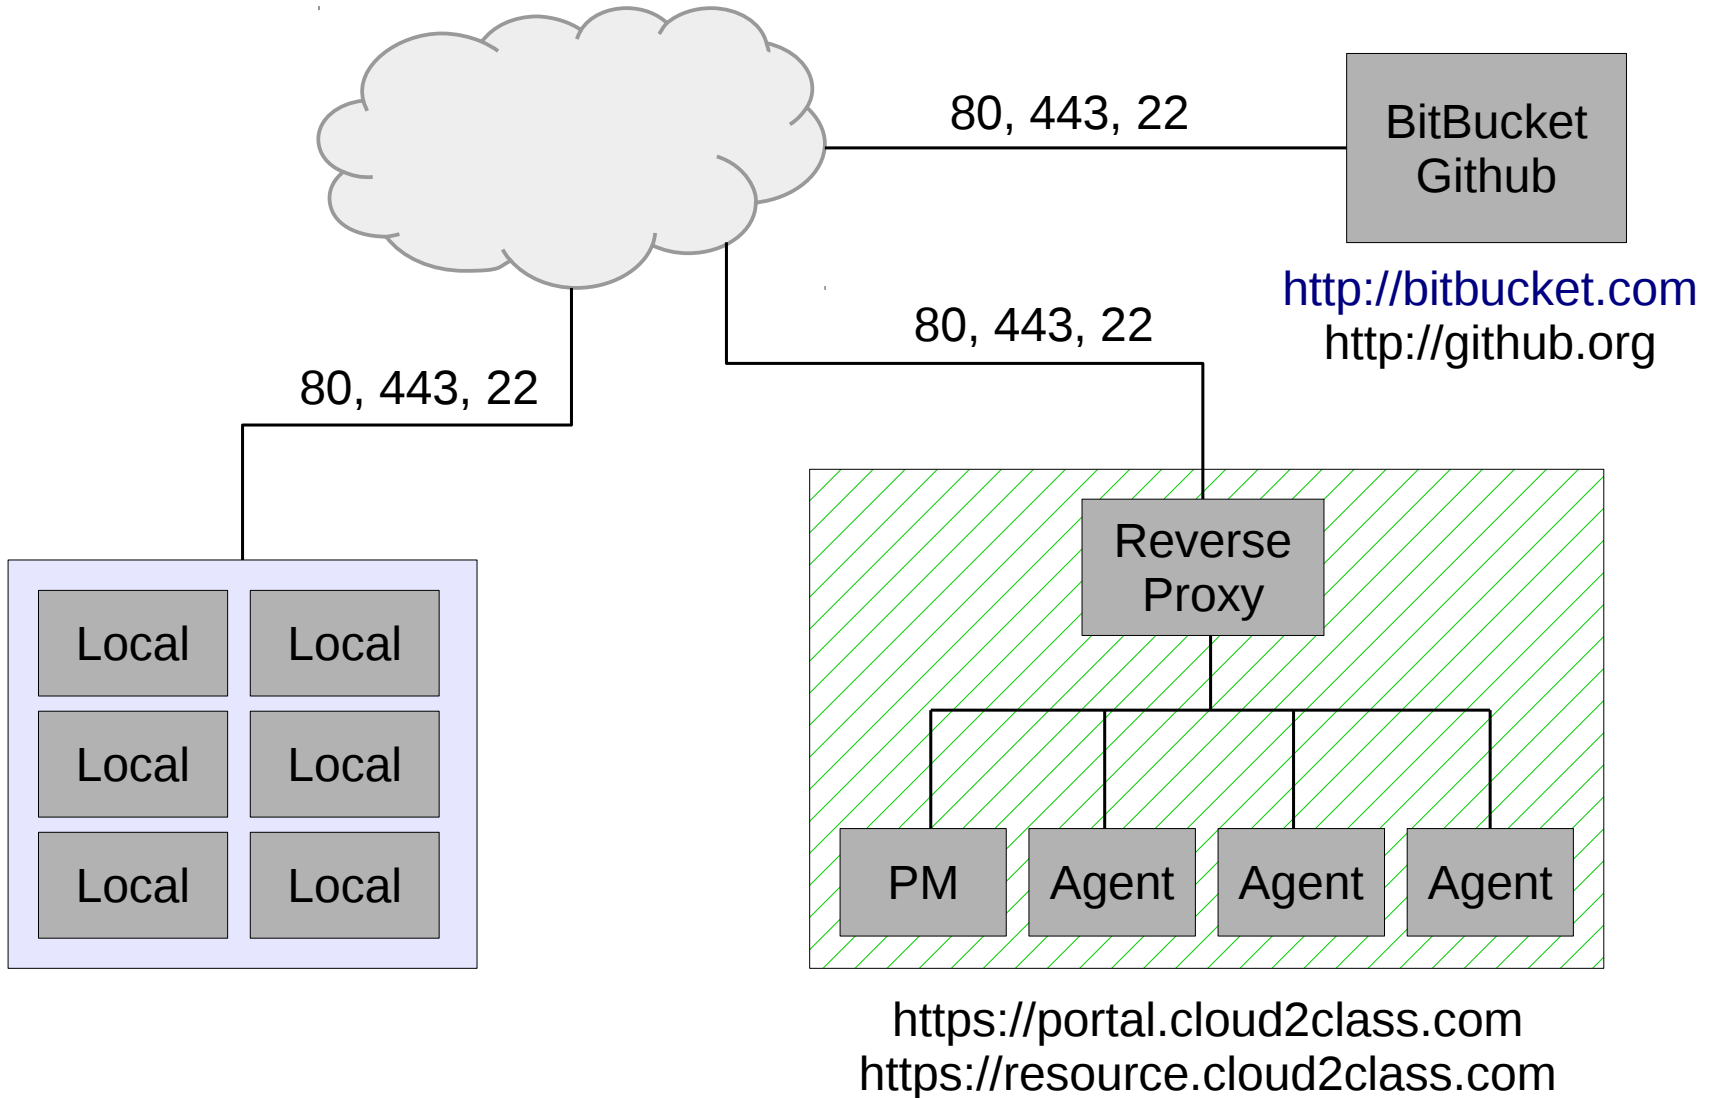
PuppetMaster instance

*2 CPU, 4MB Ram, 20GB Storage
Red Hat / CentOS or Ubuntu*

SSH access: Port 22

Web Access: Port 80 / 443

Introduction: Training Environment



Introduction: Workshops

The general aim of the workshop will be shown here

Workshops

- A set of practical tasks
- Aimed at achieving a specific outcome.
- An opportunity for discussion
- Other workshops will frequently depend on each other

If you think you cannot successfully complete a section, ask one of our instructors to help you out.

If at any point you accidentally damage your training instance, notify us and we will reset it for you

Workshop : Registration

Log on to the cloud2class platform. This will create user accounts for your instance(s), and allow access to other training resources.

Log on to the student training portal

- Browse to
- Enter your <https://portal.cloud2class.com> first and last names
- Enter the lesson security details

Security Details

*Word: *** monkey*

*Number: ***number*

Introduction: Puppet

What it IS

Configuration Management

Powerful

Widely used

Scalable to thousands of nodes

Used to define an end-state

A complete replacement for
system administration

What it is NOT

A panacea

Simple to configure

A filestore

The only option: Chef, Ansible


Capable of operating alone

A complete replacement for
system administration

Puppet: Components

Component	Function
Agent	Enacts system changes
Master	Compiles manifests for agents. Distributes your server configuration
PuppetDB	Stores node state, run logs, and facts
Puppet Dashboard	Monitor and report on your nodes
Mcollective	Orchestration tool
Facter	System information collection
Hiera	Hierarchical configuration database


Puppet: Components



Agent and Master : Index

- Agent and Master interaction
- Running the agent
- Locating the Puppet-master
- Running the Puppet-master
- Signing Certificates
- Configuration
- Webservers
- Workshop : Configuring a Puppetmaster

Agent and Master : Interaction



Agent and Master : **Running the Agent**

Puppet agent:

*Listens on port 8139
Logs to syslog
Certs are in /var/lib/puppet/ssl*

Package requirements:

*Puppet
Facter
Hier*

Starting the Puppet Agent.

```
sudo puppet agent --enable (Ubuntu)  
sudo puppet agent  
sudo service puppet start
```

Debugging / Viewing the output

```
sudo puppet agent -v -t
```

Agent and Master : Locating the puppet-master

How do we find a Puppetmaster?

- DNS - “puppet”
- DNS – srv (experimental)
- Puppet.conf

```
[agent]  
server = np-puppet.openstacklocal
```

- Command line

```
sudo service puppet agent --server
```

Agent and Master : **Running the Master**

Puppetmaster:

*Listens on port 8140
Logs to syslog
Certs are in /etc/puppet/ssl*

Package requirements:

*Puppet
Facter
Hiera
deep_merge (gem)*

Starting the Puppetmaster

```
sudo puppet master  
sudo service puppetmaster start
```

Debugging / Viewing the output

```
sudo puppet master -v -d --no-daemonize
```

Agent and Master : **Signing Certificates**

How to ensure we configure the right machines

- Autosign – Do Not Use!

```
[main]  
  autosign = true
```

- Certificate management : Puppet Cert

Node :
puppet agent -v -t --waitforcert 60

Puppet-master :
puppet cert list
puppet cert sign
puppet cert list --all

Agent and Master : Configuration

Settings in [main] are used if a more specific section doesn't set a value.

```
[main]
certname = puppetmaster01.example.com
logdir = /var/log/pe-puppet
rundir = /var/run/pe-puppet
Basemodulepath = /etc/puppet/modules
server = puppet.example.com
user = pe-puppet
group = pe-puppet
archive_files = true
archive_file_server = puppet.example.com
```

This section is used by the Puppet master and Puppet cert applications.

```
[master]
certname = puppetmaster01.example.com
dns_alt_names = puppetmaster01,puppetmaster01.example.com,puppet,puppet.example.com
ca_name = 'Puppet CA generated on puppetmaster01.example.com at 2013-08-09 19:11:11 +0000'
reports = http,puppetdb
reporturl = https://localhost:443/reports/upload
node_terminus = exec
external_nodes = /etc/puppetlabs/puppet-dashboard/external_node
ssl_client_header = SSL_CLIENT_S_DN
ssl_client_verify_header = SSL_CLIENT_VERIFY
storeconfigs_backend = puppetdb
storeconfigs = true
autosign = true
```

This section is used by the Puppet agent application.

```
[agent]
report = true
classfile = $vardir/classes.txt
localconfig = $vardir/localconfig
graph = true
pluginsync = true
environment = production
```

Agent and Master : **Webservers**

WEBrick

- Default webserver for Puppet-master
- Does not scale
- Ideal for research or proof-of-concept

Passenger

- Allows running of Rails or Rack applications within a webserver
- Scalable
- Fast

Workshop : Configuring a Puppetmaster

*Puppet has already been installed on the training machines.
Configure your machines to use themselves as the puppetmaster,
And to use Passenger as the webserver*

Configure the puppet-master

- Install the puppet-server package
- Run the puppet-master in verbose mode

Configure the client

- Configure the puppet agent to use your hostname
- Run the agent in verbose mode

Install Passenger (optional)

- Download the training repo:
`git clone https://github.com/andysingleton/training-repo.git`
- Configure puppet master to use Passenger as the webserver

Community vs Enterprise : **Index**

- Features
- Dashboard
- Support

Community vs Enterprise : **Features**

Community

- Lower Cost
- Variety of Node Classifiers (Foreman, Razor, Spacewalk)

Community

- Cloud Provisioning
- Unified Installer
- Event Inspector
- Mcollective / “Live Management”

Community vs Enterprise : Dashboard

Community

- No interaction with Mcollective
- Only partially maintained

Enterprise

- Behaves as a unit
- Better component integration
- Fully functioning dashboard
- More complicated under the hood

Community vs Enterprise : Support

Community

- Forums and User groups
- Commercial Support (**NOT** from PuppetLabs)

Enterprise

- Forums and User groups
- Commercial Support (from PuppetLabs)

Reporting : **Index**

- Configuring Reports
- Workshop : Reporting Tools

Reporting : Configuring Reports

/etc/puppet/puppet.conf

```
[master]
  reports = store , puppetdb , http
  reportdir = /var/lib/puppet/reports
  reporturl = http://np-puppet:3000/reports/upload
```

- store : yaml files are kept in “reportdir” (*/var/lib/puppet/reports/*)
- puppetdb : */etc/puppet/puppetdb.conf* determines target
- http : Pre-process them as html and forward to “*reporturl*”
- log : Use local log destinations (*syslog*)
- rrdgraph : generate RRD graphs to the “*rrddir*”
- tagmail : email logs based on tags in the messages

Workshop : Reporting Tools

*Run the puppet agent manually against the shared puppetmaster.
You should see the result on the various reporting tools*

Run puppet against another master

- Run the Puppet agent manually against np-puppet.openstacklocal
`puppet agent -v -t --server=np-puppet.openstacklocal --ssldir=/var/tmp/puppet`

View reporting tools

- PuppetDB
<https://resource.cloud2class.com/989/puppetdb>
- Puppet Dashboard (community)
<https://resource.cloud2class.com/989/dashboard>
- PuppetBoard
<https://resource.cloud2class.com/989/puppetboard>

Language : **Index**

- Overview
- Resources
- Resource Relationships
- Class Definitions
- Loading Classes
- Definition vs Declaration
- Workshop : Creating a Simple Class

Language : Overview

Resource Declarations contain :

- A Resource Type, many Attributes, Functions, Conditionals, ...

Classes contain :

- Class Definitions, Class Declarations, Resource Declarations, Defined Types, Variable assignments, Functions, Conditionals, ...

Manifests contain :

- Class Definitions, Class Declarations, Resource Declarations, Defined Types, Variable assignments, Functions, Conditionals, ...

Modules contain :

- Manifests, Templates, Files, ...

Language : Resources

- Describes an aspect of the system
- Resources are Declared – This includes them in the catalog
- The block of code is called a Resource Declaration
- Typically in Classes or Defined Types (which are then declared)

A Class containing a Resource:

```
class configure_ssh {  
  file { ["/etc/ssh/sshd_config":  
    ensure => 'present'  
    owner   => 'root',  
    group   => 'root',  
    mode    => '0600',  
    source  => 'puppet:///modules/configure_ssh/sshd_config',  
  ]  
}
```

Language : Resources

Common Resource Types

- **File:**
Covers any file, directory, link, device-file or symlink
- **Package and Service:**
Used to deploy software and manage the state of services
- **User:**
Create and manage the presence/absence of users
- **Exec**
Used to perform manual tasks: DO NOT USE IT!

<https://docs.puppetlabs.com/references/latest/type.html>

Language : Resource Relationships

Before

Require

Notify

Subscribe

Chaining Arrows (-> and ~>)

Language : Resource Relationships

before

- Apply this resource first

require

- Apply another resource first

```
class nginx {  
  package { 'nginx':  
    ...  
    before => Service['nginx'],  
  }  
  
  service { 'nginx':  
    ...  
    require => Package['nginx'],  
  }  
}  
  
Package['nginx'] -> Service['nginx']
```


Language : Resource Relationships

notify

- Tell another resource when we have been changed

subscribe

- Update our resource when another resource has changed

```
class nginx {  
  file { '/etc/nginx/nginx.conf':  
    ...  
    notify => Service['nginx'],  
  }  
  
  service { 'nginx':  
    ...  
    subscribe => File['/etc/nginx/nginx.conf'],  
  }  
}  
  
File['/etc/nginx/nginx.conf'] ~> Service['nginx']
```

Language : Class Definitions

```
class 'webusers' (  
  ensure = $ensure,  
) {  
  user { "sarah":  
    ensure      => $ensure,  
    uid         => 1001,  
    gid         => 1001,  
    firstname   => 'Sarah',  
    lastname    => 'Conner',  
    home        => $name,  
    shell       => '/bin/bash',  
    ssh-key     => 'rsa-.....',  
    managehome => true,  
  }  
}
```

Language : Loading Classes

Autoloading

- When a class or resource is declared, puppet uses its full name to find it in the available modules.
- If it can't find the class, you can't use it

Naming rules

- Base classes should be defined in `init.pp`
- Other classes should have the format `module::class`, and be defined in their own `.pp` file
- Further namespaced classes will be searched for in sub-directories and their own `.pp` file

modules/apache/manifests/init.pp

```
class apache {  
  ...  
}
```

modules/apache/manifests/vhost.pp

```
class apache::vhost {  
  ...  
}
```

modules/apache/manifests/vhost/config.pp

```
class apache::vhost::config {  
  ...  
}
```

Language : Definition vs Declaration

Defining a Class : Making it available

```
class my_own_class {  
  Resource declarations go here  
}  
  
Class my_other_class (  
  $my_parameter="my default value",  
) {  
  Resource declarations go here  
}
```

Declaring a Class : Making it happen

```
class { myclass:  
  my_parameter => "/etc/ssh/"  
}  
  
include myclass  
hiera_include ('myclass')  
contain myclass  
require myclass
```

Workshop : Creating a simple class

*Put together a simple class containing a resource definition
to deploy an executable when the class is declared
and to start a service.*

1 : Checkout the training repo

- *git clone <https://github.com/andysingleton/training-repo.git>*

2 : Create a Module

- Create a directory under */etc/puppet/modules*
(The convention is alphanumeric, and underscores)
- Create the following directories under your module directory
 - ./manifests
 - ./templates
 - ./files

Workshop : Creating a simple class

*Put together a simple class containing a resource definition
to deploy an executable when the class is declared
and to start a service.*

3 : Create a Manifest

- Create *init.pp* under “manifests” using your editor of choice

4 : Create a Class that...

- Deploys *turret.py* from files to */usr/sbin/turret.py*
- Deploys *turret-sysv* from files to */etc/init/turret*
- Starts the turret service after its files have been deployed

5 : “Classify” your node

- Add the name of your class to */etc/puppet/manifests/site.pp*
- Run puppet on your local node

Language II : **Index**

- Variables
- Scoping
- Facter
- Conditionals
- Workshop : Extending Facter

Language II : Variables

Important variable concepts :

- Assignment : `$a = 3`
- Resolution
- Interpolation : `'$variable'` vs `"$variable"`
- Scope

Language II : Variables

Boolean :

- `$myvar = true`

String :

- `$myvar = 'somevalue'`

List :

- `$myvar = ['first' , 'second']`

Hash :

- `$myvar = { 'name' => 'consuela' }`

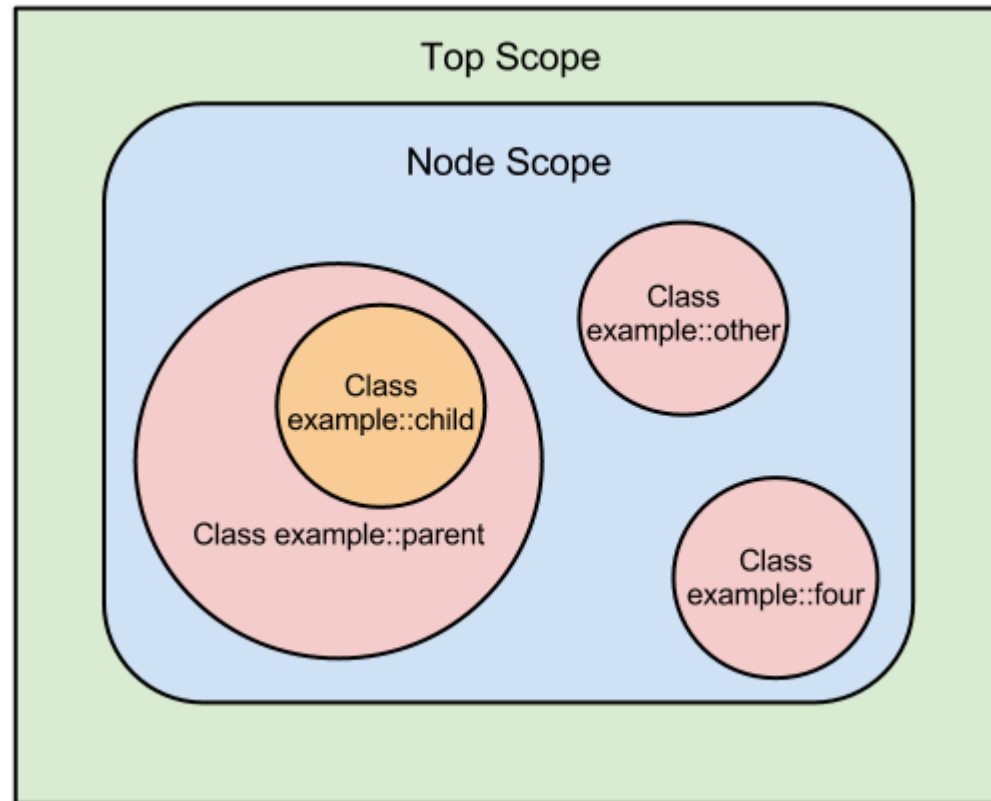
Nested :

- `$myvar = { ['name' => 'consuela' , 'shell' => '/bin/bash'] }`

```
class manage_ssh (
  $permitrootlogin = 'no',
  $ensure = 'running',
  $enable = true,
){
  case $osfamily {
    'Debian': { $servicename = 'ssh' }
    'RedHat': { $servicename='sshd' }
  }
  service { $servicename:
    ensure    => $ensure,
    enable    => $enable,
    hasrestart => true,
    hasstatus => true,
    require   => Package['openssh-server']
  }
}
```

Language II : Scoping

“Variables in a given scope are only available within that scope and its children, and any local scope can override the variables it receives from its parents.”



Language II : Scoping

Top scope

- Variables defined outside of a class, type, or node definition at top scope. e.g. in site.pp

Node scope

- Variables defined inside a node definition are available everywhere for that specific node

Local scope

- Variables defined inside a class or defined type are in local scope.
- Available only within the scope and its children.

```
/etc/puppet/manifests/site.pp
```

```
$purple = "top scope"  
$green = ["top scope list"]
```

```
node 'mynode.example.com' {  
  $yellow = 'node local scope'  
}
```

```
class core {  
  $blue = "local scope"  
  $green += ['appended data']  
}
```

```
class a_n_other {  
  $blue = "another local scope"  
  $purple = "overridden local scope"  
}
```

Language II : Scoping

Global Qualified Name

- A method to access any variable anywhere:
\$apache::params::confdir

Language II : **Facter**

Core Facts

- Node information always available from facter

Custom Facts

- Written in Ruby
- Distributed manually or from a module (*[module]/lib/facter/[fact].rb*)
- Loaded in puppet.conf using '*pluginsync = true*', and '*facter -p*'

External Facts

- Arbitrary executables, or static files (e.g. yaml)
- Distributed from a module (*[module]/facts.d/[your binary]*)
- Loaded in puppet.conf using '*pluginsync = true*', and '*facter -p*'
- Manually drop into */etc/facter/facts.d/*

Language II : **Facter**

- Facter collects system information in the form of “Facts”.

facter -p

- Facts can be resolved in the same manner as any variable

```
if $fact_name == 'ok' {  
  ensure => $facts['fact_name']    *requires trusted_node_data = true
```

- Can also be referenced as `$::fact_name` for clarity

Language II : Conditionals

if / elsif / else

- If expression evaluates to true, execute code block

unless

- If expression evaluates to false, execute code block

case

- Execute code block relating to matched value

selector

- Return a value instead of executing a code block

```
class deploy_ssh {  
  if $osfamily == 'Debian' {  
    $ssh_service = 'ssh'  
  }  
  elsif $osfamily == 'RedHat' {  
    $ssh_service = 'sshd'  
  }  
  else {  
    $ssh_service = 'ssh'  
  }  
  
  case $osfamily {  
    'Debian': { $ssh_service = 'ssh' }  
    'RedHat': { $ssh_service = 'sshd' }  
    default: { $ssh_service = 'ssh', }  
  }  
  
  $ssh_service = $osfamily ? {  
    'Debian' => 'ssh',  
    'RedHat' => 'sshd',  
    Default => 'ssh',  
  }  
}
```

Workshop : Extending Facter

We are going to create an external fact that checks the current state of our Turret service, then add some conditional logic based on it.

1 : Create an external executable fact

- Either use the training-repo “*check_turret.sh*” script or create your own; We want to know if the service is running or not
- Check “*facter*” responds correctly

2 : Populate `/var/www/html/index.html` with a status page

- Extend your class to change the content of `index.html`
- If the turret service is up, provide a web-page stating that everything is fine
- If the service is down, provide a page saying the service is down

Language III : Index

- Templates
- Workshop : Creating a Template
- Defined Types
- Workshop : Creating a Defined Type

Language III : Templates

Templates are an easy way to distribute similar files to different nodes.

What is a template :

- Plain text plus embedded Ruby code
- A “.erb” file in *[module]/templates*

How to use it :

- With the “template” function from within a File resource
- Some knowledge of Ruby is good

Benefits :

- Can contain complex logic
- Keeps the filestore clean
- As simple or as complex as you want them to be

Language III : Templates

Ruby Expression

- Replace the content of the tags with the output of the code
- This can include variables from your manifests

PermitRootLogin <%= @permitroot -%>

Ruby Code

- Execute any code. It will not be replaced by any value in the final content.

PermitRootLogin <% if @mycondition == false -%>

Comments

- Suppressed in the final output

<%# We include this loop to workaround bug #659 %>

<%-

- Suppress leading whitespace

-%>

- Suppress following linebreak

Language III : Templates

Scoped Variables :

- All variables in your current scope can be used with an @
PermitRootLogin <%= @permitroot -%>

Out of scope variables :

- Any variable can be referenced with the `scope.lookupvar` method
<%= scope.lookupvar('apache::user') %>

Language III : Templates

modules/ntp/manifests/config.pp (snippet)

```
file { $config:
  ensure => file,
  owner  => 0,
  group  => 0,
  mode   => '0644',
  content => template($config_template),
}
```

modules/templates/authorized_keys.erb

```
<% [@servers].flatten.each do |server| -%>
server <%= server %><% if @preferred_servers.include?(server) -%> prefer<% end %>
<% end -%>
```

```
<% if scope.lookupvar('::is_virtual') == "false" or @udlc -%>
# Undisciplined Local Clock. This is a fake driver intended for backup
# and when no outside source of synchronized time is available.
server 127.127.1.0
fudge 127.127.1.0 stratum 10
restrict 127.127.1.0
<% end -%>
```

Language III : Templates

modules/usermgmt/manifests/user.pp (snippet)

```
file { "${home}/.ssh/authorized_keys":  
  ensure => $ensure,  
  owner  => $name,  
  group  => $name,  
  mode   => '0600',  
  content => template("${module_name}/authorized_keys.erb"),  
  require => File["${home}/.ssh"],  
}
```

modules/templates/authorized_keys.erb

```
<% @ssh_pub_keys.sort.each do |key| -%>  
<%= key %>  
<% end -%>
```

Language III : Templates

modules/usermgmt/manifests/user.pp (snippet)

```
file { "${home}/.ssh/authorized_keys":  
  ensure => $ensure,  
  owner  => $name,  
  group  => $name,  
  mode   => '0600',  
  content => template("${module_name}/authorized_keys.erb"),  
  require => File["${home}/.ssh"],  
}
```

modules/templates/authorized_keys.erb

```
<% @ssh_pub_keys.sort.each do |key| -%>  
<%= key %>  
<% end -%>
```

Language : **Defined Types**

Defined Types are repeatable pieces of configuration
They are composed of other resources, and behave
exactly as any other resource.

Creating a Defined Type

- The define keyword
- A name
- A list of parameters
- A block of code

Language : Defined Types

```
define planfile ($user = $title, $content) {  
  file {["/home/${user}/.plan":  
    ensure => file,  
    content => $content,  
    mode   => 0644,  
    owner  => $user,  
    require => User[$user],  
  }  
}  
  
user {'nick':  
  ensure   => present,  
  managehome => true,  
  uid      => 517,  
}  
planfile {'nick':  
  content => "Some new content",  
}
```

Language : **Defined Types**

Defined Types are repeatable pieces of configuration
They are composed of other resources, and behave
exactly as any other resource.

Creating a Defined Type

- The resources inside the type must depend on parameters